


Figure 1. The multilingual token tax across 16 languages × 3 tokenizers



Tokens to encode the same parallel FLORES-200 sentence, ÷ English. Brahmic-derived & mainland-SEA scripts pay 7–11× on the legacy GPT-4/3.5 tokenizer (Burmese worst, 11.2×); the GPT-4o tokenizer cuts most to ~2×.