


Figure 2. Context-window exhaustion rate by language (GPT-4/3.5 tokenizer)



Deterministic: cumulative tokens = (100 words/turn) × fertility × turn. Tamil/Malayalam fill the 4096 window in ~3 turns; English takes ~32. Not an accuracy claim.